

Parallel programming

// programming

S.I

Françoise Baude

Université Côte d'Azur

Polytech Nice Sophia

baude@unice.fr

web site: <https://lms.univ-cotedazur.fr/course/view.php?id=14214>

(self registration)

Feb 2023

Chapter 2: PRAM related measures, Complexity

Plan

1. Some indicators to evaluate a PRAM algorithm
2. Complexity of parallel problems

Work of PRAM algorithms

- SURFACE : number of processors used by a PRAM algorithm $A_{//}(N)$ working on a size N problem:
 - $H(A_{//}(N))$ = the maximum amount of procs required in a given parallel PRAM instruction, during the algo
- PARALLEL TIME $T_p(A_{//}(N))$ of the algo using P procs:
 - $T_p(A_{//}(N))$ = the number of computation steps when using P procs
- WORK : product of SURFACE by PARALLEL TIME
 - $$W = H(A_{//}(N)) * T_{H(A_{//}(N))}(A_{//}(N))$$
$$= P * T_p$$

Speedup, Efficiency

- Consider the (best)Seq time to solve the problem $A_{seq}(N)$ whose time is $T_{seq}(N)$ (using 1 proc!)
- SPEEDUP (acceleration factor) of $A_{//}(N)$ using P procs
 - $S_p(N) = T_{seq}(N) / T_p(A_{//}(N))$
 - Theoretical goal is to have $S_p(N) = p$
- EFFICIENCY of $A_{//}(N)$ using P procs
 - $e = \text{Sequential work} / \text{Parallel work}$
 - $e = T_{seq}(N) * 1 / W$
 $= T_{seq}(N) / p * T_p(A_{//}(N))$
 - Theoretical goal is to have $e=1$


By using p procs,
the speed-up is p

Seq runtime has
been converted into
adding procs!

Speedup curves

Speed-up is a factor, not a speed, nor a duration

Speedup curves measure the utility of parallel computing, not speed.



trivially parallel

(e.g., matrix product, LU decomposition, ray tracing)
→ close to ideal $S = p$

work-bound algorithms

→ linear $SU \in \Theta(p)$, work-optimal

tree-like task graphs

(e.g., global sum / max)
→ sublinear $SU \in \Theta(p / \log p)$

There is a high variation in the number of proc use during the computation

communication-bound

→ sublinear $SU = 1 / fn(p)$

Most papers on parallelization show only relative speedup

(as $SU_{abs} \leq SU_{rel}$, and best seq. algorithm would be needed for getting Su_{abs})

Speed-up in practice : ways to measure performances

- Theoretical Speed-up (Absolute speedup):
 - Computed using the complexity of the algorithm solving problem in sequential
- Can we always have the sequential time ?
 - Sometimes, no sequential implementation exists
 - Take the parallel version, run it using $p = 1$
 - Sometimes, not feasible when n is too big
 - Not enough memory to run with input of size n
 - Take the parallel version and measure its time, when p increases
 - ➔ In these cases, we measure the RELATIVE speed-up
 - Sometimes, the sequential time gets penalized due to some memory cache effects
 - With more processors, the measured speedup becomes better than the theoretical speedup ... ☺ : because mem accesses apply more often in cache
 - Speed-up becomes super-linear !

Example: Cost-optimal parallel sum algorithm on SB-PRAM

Saarbrücken Univ. PRAM: real PRAM machine [1990]

$n = 10,000$

Processors	Clock cycles	Time	SU _{rel}	SU _{abs}	EF _{rel}
Sequential	460118	1.84			
1	1621738	6.49	1.00	0.28	1.00
4	408622	1.63	3.97	1.13	0.99
16	105682	0.42	15.35	4.35	0.96
64	29950	0.12	54.15	15.36	0.85
256	10996	0.04	147.48	41.84	0.58
1024	6460	0.03	251.04	71.23	0.25

$$\frac{T_{\text{par}}(1)}{p \cdot T_{\text{par}}(p)}$$

$n = 100,000$

Processors	Clock cycles	Time	SU _{rel}	SU _{abs}	EF _{rel}
Sequential	4600118	18.40			
1	16202152	64.81	1.00	0.28	1.00
4	4054528	16.22	4.00	1.13	1.00
16	1017844	4.07	15.92	4.52	0.99
64	258874	1.04	62.59	17.77	0.98
256	69172	0.28	234.23	66.50	0.91
1024	21868	0.09	740.91	210.36	0.72

Theorem: *Conservation of work by simulation*

- Given an algorithm A running in time t on p procs. of a given PRAM


*It is possible to simulate A on a same sort of PRAM using $p' \leq p$ procs in time $O(t * (p / p'))$*

- Proof: intuitive! At each step, each p' proc will have to execute a subset of (p/p') // instructions in sequence
- Work is kept as it is:
 - $W = p * t$, new $W = p' * (t * (p/p')) = t * p$

Exemple: EREW maximum computation -v2 => v3

$m=k=8, k'=6$

$N=2^k=2^8 = 256$. $p=2^{(k-1)}=2^7=128$. $p'=2^{(k'-1)}=2^5=32$
each of the $O(\log(2^8))=O(8)$ instructions of the
 $\text{Maxv2}_{//}(2^8)$ will be simulated by up to 2^5 procs,
executing up to $2^{(k-k')}=2^2=4$ max binary operations



Pour ($k=m-1; k \geq 0; k--$) **/*still costs m parallel steps*/**

/*enroll up to $2^{(k-1)}$ procs per step ? NO, just enroll up to $q=2^{(k'-1)}$ */

Pour chaque proc q en parallele

Pour ($l=0; l < 2^{(k-k')}; l++$) **/*costs $2^{(k-k')}$ seq time*/**

$A[zz] = \max(A[yy], A[yy+1]);$

Consequences of Theorem:

Conservation of work by simulation

- The work of a $A_{//}(N)$ with p proc is **at least** in the order of $T_{seq}(N)$, the best sequential algo


$$H(A_{//}(N)) * T_p(A_{//}(N)) \geq T_{seq}(N)$$

$$T_{seq}(N) \leq H(A_{//}(N)) * T_p(A_{//}(N))$$

- Proof: by absurd
 - Suppose $H(A_{//}(N)) * T_p(A_{//}(N)) < T_{seq}(N)$,
 - then, decide $p'=1$,
 - Each p' proc will have to execute a subset of (p/p') // instructions in sequence
 - If $T_1(A_{//}(N)) < T_{seq}(N)$, the seq algo was not the best !!

WORK EFFICIENCY

- A parallel algorithm is said to be WORK EFFICIENT:
 - Its work is of the same amount than the **best sequential** algorithm
 - i.e. $e == 1$
- (counter-)Examples:
 - *max-v3* versus $O(n)$ seq max.
 - Using the simulation theorem
 - $O(t * (p / p'))$; $p=n$; choose $p'=n/\log n$; $t = \log n$
 - *max-v3* time = $\log n * n/n/\log n = \log^2 n$ (limiting factor: same t)
 - *max-v3* work = $n/\log n * \log^2 n = n * \log n \Rightarrow$ not work efficient
 - Because of the same $t=\log n$ on same initial size= n
 - *max-v2 with subtrees*, versus $O(n)$ seq max.
 - $[\log n + O(\log (n/\log n))] * (n/\log n) = O(n)$, work eff.
 - Here *max-v2 with subtrees* applies the Brent Principle



Brent Principle

- A general principle to decrease number of used procs (not a method, just a principle!)
- Given $A_{//}(N)$ having a total of m operations, running in $t = T(A_{//}(N))$ with an unbounded number of procs

One can simulate $A_{//}(N)$ in time $O(m/p + t)$ on a similar PRAM using p processors

- Proof: at step i , $A_{//}$ runs $m(i)$ ops s.t. $\sum m(i) = m$
 - Simulation with p procs takes $\lceil m(i)/p \rceil \leq m(i)/p + 1$

In total, the time of simulation $\leq \sum_{i=1 \text{ to } t} (m(i)/p + 1) = m/p + t$


- Consequence: if $(m/p) = t$ then $T_p(A_{//}(N)) \leq 2 \cdot t$

(Proof of) Brent has no real incidence!

- The proof is not constructive
 - It is not a method, it is not a recipe in order to go from an unlimited number of procs to a practical & concrete number of p processors
 - It does not tell how to concretely split m operations into $m(i)$
 - For a given step i , it does not tell how to split the $m(i)$ operations into p tasks (it is not a load sharing method)
 - E.g list ranking on a linked list of size n : it is not easy to split the list into p sublists, each of successive elements, of size n/p
 - Still, in some cases, like working on linear structures such as arrays, quite easy to be split into seq tasks then parallel tasks, eg:

■ Ex: on max-v2, how many max ops ? $\sum m(i) = m$

- $k=8, 2^7=128$; + $k=7, 2^6=64$; + $2^5=32$; ... + $k=0, 2^0$
- Choose $p'=p/t=n/\log n=256/8=32$; $t=8 \Rightarrow t'=\log(n/\log n)=5$
- New Step1: $256/32=2^8/2^5=8$ values to work on, in seq. per proc, $t_{seq}=8 \Rightarrow$ we needed to modify the // algo (8x32 max ops)
- Next steps: just parallel max-v2 working on 32 values: $t'=5$



Plan

1. Some indicators to evaluate a PRAM algorithm
2. **Complexity of parallel problems**

The NC complexity class of parallel problems


- It tell us what is a « good » parallel algo
- NC complexity class (« Nick's Class »)
 - The set of problems for which there exists a parallel algorithm taking a (poly)logarithmic parallel time, and using a polynomial number of processors
 - $NC \text{ in } \mathcal{P}$, $\mathcal{P} \text{ in } NC$, probably \neq
- An « Optimal » parallel algorithm :
 - Belongs to NC, and moreover is efficient (in work)
- Be careful in practice with the poly-log time:
 - Ex: $A // \text{time} = \log^3 n \ll n^{1/4}$ only when $n > 10^{12}$
- More: [Parallel complexity theory - NC algorithms \(wisc.edu\)](http://wisc.edu) and [Parallel complexity theory - P-completeness \(wisc.edu\)](http://wisc.edu)

Sort n values in an optimal way ?


- Complexity to sort n values:
 - Somehow, they all must be compared 2 by 2 (cf max-v1)
 - It is known that the lower bound in sequential is $\Omega(n \cdot \log n)$
 - => It provides us with a framework for parallel algorithms!
 - **With only $O(n)$ procs. used, goal is to sort in $O(\log n)$ //time**
 - It is feasible, but the factor hidden in the $O()$ is very high
 - Principle of the merge parallel sort algo on an EREW [due to Cole] :
 - Start from n lists of size 1, merge them two by two in //
 - Start again, to merge all these lists 2 by 2, and so on
 - Depth of the tree to traverse from leaves to root: $\log n$, so, the sub goal is to **merge two lists of length $O(n)$ in constant time** It is hard but feasible ; make inactive procs of the upper stages in the tree become active in order to contribute to these merge operations in $O(1)$ // time that run at lower stages: pipeline, anticipate
- In practice: Sort in // in time $O(\log^2 n)$, non optimal
 - On a PRAM
 - Or, on a « sorting network » : it is a topology of *sorting elements* that is always the same whatever be the initial sequence of input data to be sorted
 - Provides an « oblivious » or « regular » algorithm (i.e., just dependant of the problem size, not of the data values)

Architecture of a sorting network


- Built using 2x2 comparators/sorting elements
- Architecture of comparison-exchange sorting networks.
 - (a) The default type of comparator
 - (b) The second type of comparator.
 - (c) Sorting network composed from columns of basic comparators.



(a)

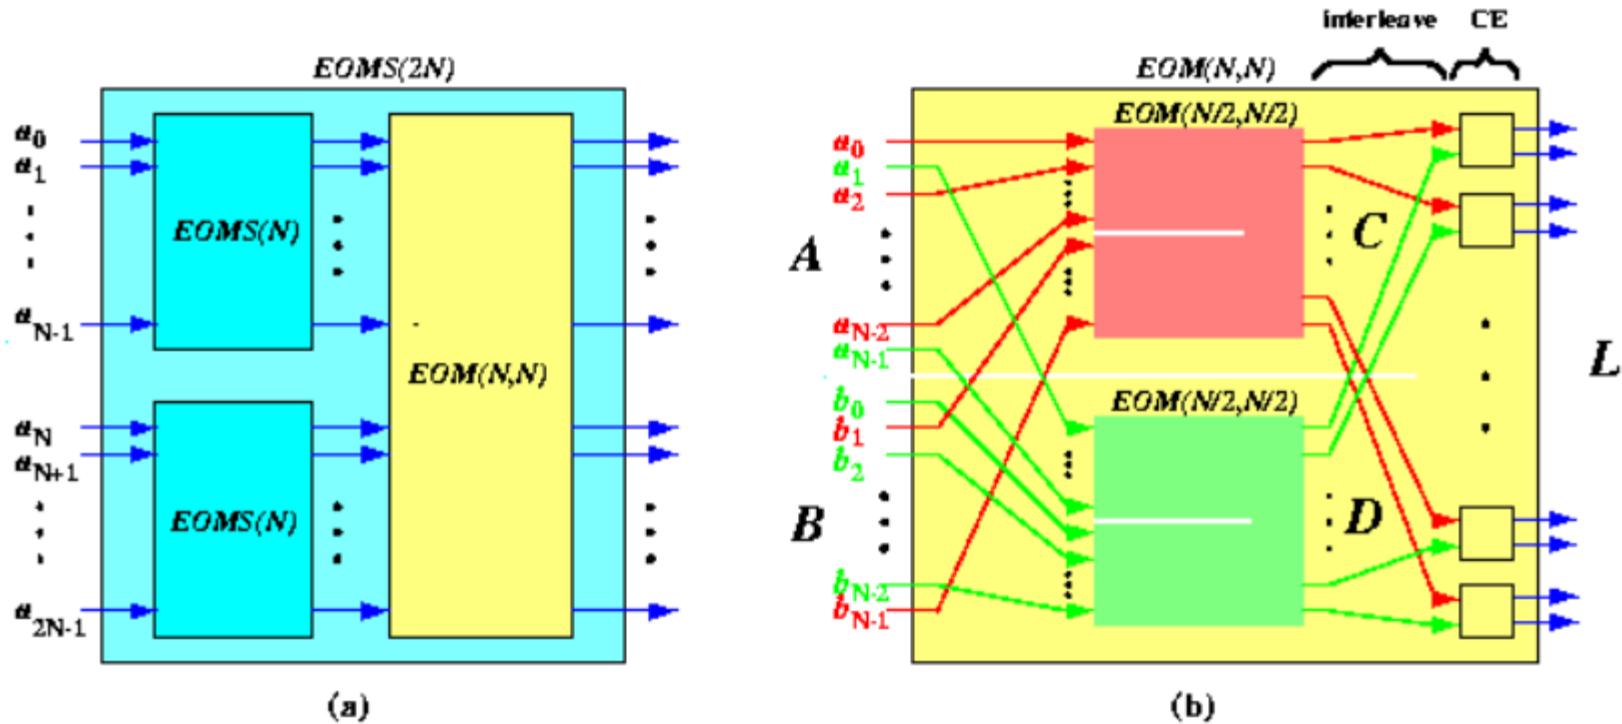


(b)



(c)

Even-Odd Merge Sorting network



Even-Odd MergeSort (a) and Merge (b) network

$$C = \{EOMerge\}(even(A), odd(B))$$

$$D = \{EOMerge\}(odd(A), even(B))$$

$$L' = \{Interleave\}(C, D)$$

$$L = \{EOMerge\}(A, B) = \{Pairwise_CE\}(L')$$